

Introduction to Error-Correcting Codes

Pierre Fromonot

October 2, 2026

Contents

1	Theory of Error-Correcting Codes	2
1.1	Motivation	2
1.2	Codes and Linear Codes	2
1.2.1	Codes	2
1.2.2	Linear Codes	2
1.2.3	Notation	3
1.3	Generator and Parity-Check Matrices	3
1.4	Syndrome Decoding	3
1.5	Distance and Error Correction Capability	4
1.5.1	The Minimum Distance	4
1.5.2	The Error-Correction Bound	4
1.5.3	Geometric Intuition	4
1.5.4	Decoding Algorithm	5
1.5.5	Code Parameters and Trade-offs	5
1.6	Systematic Forms of Codes	5
1.6.1	Intuition: Information and Redundancy	6
1.6.2	Systematic Forms Always Exist	6
1.6.3	Relationship with the Parity-Check Matrix	6
1.7	Bounds on Code Parameters	7
2	Examples of Codes	9
2.1	Hamming Codes	9
2.1.1	The idea	9
2.1.2	Example: the $[7, 4]$ Hamming code	10
2.1.3	SECDED Codes	10
2.2	Reed-Solomon Codes	11
2.3	Low Density Parity Check (LDPC) Codes	11
3	Error-Correcting Codes in Cryptography	12
3.1	Motivation	12

Chapter 1

Theory of Error-Correcting Codes

1.1 Motivation

Error-correcting codes are mathematical objects that, combined with the right algorithms, detect and correct data errors by adding redundancy to transmitted or stored information. This allows systems to recover original data without requiring re-transmission, which is critical in noisy channels, memory storage, and deep-space communications.

1.2 Codes and Linear Codes

1.2.1 Codes

At its core, a *code* is simply a set of allowed sequences (called *codewords*) that we agree upon before communication. The sender chooses a codeword from this set, transmits it through a noisy channel, and the receiver tries to decode it, correcting any errors introduced by noise.

Definition 1.2.1. A (binary) code of length n is a non-empty subset $C \subseteq \mathbb{F}_2^n$. Each element of C is called a codeword. The rate of the code is $\log_2 |C|/n$.

Example 1.2.2. A simple example is the repetition code: to send one bit, repeat it three times. Here $C = \{000, 111\} \subset \mathbb{F}_2^3$. If a single error occurs (e.g., $000 \rightarrow 010$), we can still decode by majority vote.

1.2.2 Linear Codes

In principle, one could study any subset of $\mathbb{F}_2^n$. However, linear codes — codes that form a vector subspace — have special structure that makes them practical:

- **Algebraic structure:** Linear operations (addition, scalar multiplication) are well-defined, allowing us to leverage tools from linear algebra.
- **Efficient encoding:** We can represent the code compactly using a generator matrix G , which encodes any message via a simple matrix-vector product.
- **Efficient decoding:** The structure of linear codes enables fast decoding algorithms, such as syndrome decoding (which we will discuss later).
- **Distance properties:** For linear codes, the Hamming distance between codewords is directly related to the weight of other codewords, simplifying analysis.

Fortunately, many of the best codes discovered are linear. For this reason, we will focus exclusively on linear codes.

1.2.3 Notation

We denote a linear code with parameters n , k , and minimum distance d as an $[n, k, d]$ code, where:

- n is the **length** (number of bits in a codeword),
- k is the **dimension** (log-number of codewords: $|C| = 2^k$),
- d is the **minimum distance** (smallest Hamming distance between distinct codewords; we define this formally later).

1.3 Generator and Parity-Check Matrices

Formally, a linear code C is a vector subspace of $\mathbb{F}_2^n$ of dimension $k \leq n$. We write G a *generating matrix* ($n \times k$) of C : its rows form a basis of C ,

$$C = \{x \cdot G : x \in \mathbb{F}_2^k\}$$

Here x is a *message*, and $c = x \cdot G \in \mathbb{F}_2^n$ is a *codeword*. As $n \geq k$, computing $c = x \cdot G$ adds redundancy to the message x .

When transmitted, $c = x \cdot G$ may be corrupted by an error which can be seen as a vector $e \in \mathbb{F}_2^n$: $y = c + e = x \cdot G + e$. The redundancy added to x can help recover the actual message x in some cases when the code C is well structured and if the error e is not too large (if the error occurs on $t \ll n$ bits, i.e., e has a small Hamming weight).

We define the *dual code* of C as

$$C^\perp = \{v \in \mathbb{F}_2^n : v \cdot G = 0\}$$

This is the left kernel of G . The *parity matrix* H ($m \times n$) of the code C is defined by

$$H \cdot c^T = 0, \quad \forall c \in C$$

where $m = n - k$ is the number of *parity bits*. This is a generating matrix of $C^\perp$.

1.4 Syndrome Decoding

Here is why H is important. When a codeword $c \in C$ is transmitted through a noisy channel, the received word is $y = c + e$, where e is the error vector. Although the receiver does not know c or e directly, it can compute the *syndrome*:

$$s = Hy^T = H(c + e)^T = Hc^T + He^T = He^T,$$

since $Hc^T = 0$ for all $c \in C$ (by definition of the parity-check matrix).

The syndrome s is a crucial quantity: it depends only on the error e , not on the message or codeword. We can think of the syndrome as a signature of the error. Given the syndrome, the decoder's goal is to determine which error e most likely occurred, and thus recover the original codeword as $c = y - e$.

For this strategy to work, we need the error to be small enough. In particular, if two distinct errors e_1 and e_2 produce the same syndrome, then the decoder cannot distinguish between them, and decoding fails. This motivates the concept of the minimum distance, which tells us the maximum number of errors that can always be corrected uniquely.

1.5 Distance and Error Correction Capability

1.5.1 The Minimum Distance

For a linear code C , the *minimum distance* is defined as the smallest Hamming distance between two distinct codewords:

$$d := \min_{\substack{c_1, c_2 \in C \\ c_1 \neq c_2}} d_H(c_1, c_2).$$

A key property of linear codes is that this is equivalent to the minimum weight of non-zero codewords:

$$d = \min_{c \in C \setminus \{0\}} w_H(c).$$

This holds because $d_H(c_1, c_2) = w_H(c_1 - c_2) = w_H(c_1 + c_2)$ (in $\mathbb{F}_2$), and $c_1 + c_2 \in C$ whenever $c_1, c_2 \in C$.

1.5.2 The Error-Correction Bound

The minimum distance directly determines how many errors a code can correct. The intuition is geometric: if we place a ball of radius r around each codeword, and these balls are disjoint, then any received word within distance r of a codeword can be uniquely decoded to that codeword.

Theorem 1.5.1. *Let C be a linear $[n, k, d]$ code. Then C can correct up to*

$$t = \left\lfloor \frac{d-1}{2} \right\rfloor$$

errors. That is, if e is an error vector with $w_H(e) \leq t$ and $c \in C$ is the transmitted codeword, then c is the unique codeword closest to $y = c + e$.

Proof. Let $y = c + e$ be the received word, where c is the transmitted codeword and e is an error of weight $w_H(e) \leq t$. We need to show that c is the unique codeword closest to y . The distance from y to any codeword c' is

$$d_H(y, c') = d_H(c + e, c') = w_H(c + e + c').$$

If $d_H(y, c') \leq t$, then $w_H(c + e + c') \leq t$. Now, $c + c'$ is a codeword (as the difference of two codewords), and let's write $c + c' =: \tilde{c}$. Then

$$w_H(\tilde{c} + e) \leq t.$$

By the triangle inequality for Hamming weight,

$$w_H(\tilde{c}) \leq w_H(\tilde{c} + e) + w_H(e) \leq t + t = 2t.$$

But $\tilde{c} = c + c'$ is a non-zero codeword (since $c \neq c'$), so $w_H(\tilde{c}) \geq d$ by definition of minimum distance. Therefore, $d \leq 2t$. However, $t = \lfloor (d-1)/2 \rfloor$, which means $t \leq (d-1)/2$, so $2t \leq d-1 < d$. This is a contradiction.

Thus, there cannot exist another codeword $c' \neq c$ with $d_H(y, c') \leq t$, so c is the unique closest codeword to y . $\square$

1.5.3 Geometric Intuition

The theorem has a natural geometric interpretation. Imagine placing a ball of radius t around each codeword in $\mathbb{F}_2^n$. The theorem says these balls are disjoint. If an error of weight at most t is introduced, the received word falls inside the ball around the true codeword, ensuring unique decoding.

More precisely, if the minimum distance is d , then the balls of radius $t = \lfloor (d-1)/2 \rfloor$ around each codeword do not overlap. To see why: if two balls around c and c' overlapped at some point y , then $d_H(c, c') \leq d_H(c, y) + d_H(y, c') \leq t + t = 2t < d$, contradicting the definition of minimum distance.

1.5.4 Decoding Algorithm

In practice, syndrome decoding proceeds as follows:

1. Receive $y = c + e$.
2. Compute the syndrome $s = Hy^T = He^T$.
3. Use a lookup table (or a decoding algorithm) to find the error vector e of minimum weight with the given syndrome.
4. Recover the codeword as $c = y + e$ and extract the message.

The key insight is that the syndrome uniquely identifies the error (assuming the error weight does not exceed t), and the decoder does not need to know the transmitted codeword or message beforehand.

1.5.5 Code Parameters and Trade-offs

A linear code is denoted $[n, k, d]$ with the following parameters:

- n is the codeword length.
- k is the dimension, so there are 2^k codewords.
- d is the minimum distance, determining the error-correction capability $t = \lfloor (d - 1)/2 \rfloor$.

A good code achieves a balance: it should have a high rate k/n (to minimize redundancy), a large minimum distance d (to correct many errors), and short codeword length n (for practical efficiency). In reality, designing codes is often a trade-off between these competing goals.




Figure 1.1: Unique decoding sphere visualization: if the minimum distance is d , then balls of radius $t = \lfloor (d - 1)/2 \rfloor$ around each codeword are disjoint, ensuring unique decoding.

1.6 Systematic Forms of Codes

In practice, it is often useful to organize the generator matrix G in a *systematic form*, where the information from the message is directly visible in the codeword.

Definition 1.6.1. A generator matrix G is in systematic form if it can be written as

$$G = (I_k \mid P),$$

where I_k is the $k \times k$ identity matrix and P is a $k \times (n - k)$ matrix of parity bits. In other words,

$$G = \begin{pmatrix} 1 & 0 & \cdots & 0 & p_{1,1} & \cdots & p_{1,n-k} \\ 0 & 1 & \cdots & 0 & p_{2,1} & \cdots & p_{2,n-k} \\ \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 & p_{k,1} & \cdots & p_{k,n-k} \end{pmatrix}.$$

1.6.1 Intuition: Information and Redundancy

When G is in systematic form, the encoding process becomes transparent. If the message is $m = (m_1, \dots, m_k)$, then the codeword is

$$c = m \cdot G = (m_1, \dots, m_k, p_1, \dots, p_{n-k}),$$

where the first k positions are exactly the message, and the last $n - k$ positions are the parity bits computed as linear combinations of m .

Why is this useful?

- The information is directly embedded in the first k positions of the codeword — no need to "decode" the message from a complicated codeword structure.
- The parity bits (the matrix P part) are redundancy: they duplicate and check the information, allowing error detection and correction.
- At the receiver, after decoding, one simply extracts the first k bits to recover the original message.

1.6.2 Systematic Forms Always Exist

The key observation is that any linear code has a generator matrix, and any generator matrix can be transformed into systematic form using Gaussian elimination.

Proposition 1.6.2. *Any linear $[n, k]$ code has a generator matrix in systematic form.*

Proof. Let G be a $k \times n$ generator matrix for a linear code C . The rows of G are linearly independent (since they form a basis of C). Perform Gaussian elimination on G to reduce it to row-echelon form. Since the rows are linearly independent and there are k rows, we can reduce G to a form where:

1. The first k columns form an identity matrix I_k (the leading entries of each row are in different positions, forming a pivot structure).
2. The remaining columns contain arbitrary values.

By permuting columns if necessary (reordering the positions of the codeword), we can arrange the matrix so that the identity appears in the first k columns. This gives us $G = (I_k \mid P)$ in systematic form. $\square$

Note: permuting columns changes the positions of information and parity in the codeword, but does not change the code itself or its properties.

1.6.3 Relationship with the Parity-Check Matrix

If $G = (I_k \mid P)$ is in systematic form, then a parity-check matrix in systematic form is

$$H = (-P^T \mid I_{n-k}) = (P^T \mid I_{n-k})$$

(over $\mathbb{F}_2$, where $-P = P$). This shows the duality between generator and parity-check matrices.

1.7 Bounds on Code Parameters

Upper bounds define the impossible frontier: they tell us what no code can achieve.

Theorem 1.7.1 (Singleton Bound). *For any $[n, k, d]$ code over any alphabet:*

$$d \leq n - k + 1.$$

Proof. Let C be a linear $[n, k]$ code with minimum distance d . Puncture (remove) any $d - 1$ coordinates. The resulting code C' has length $n - (d - 1)$ and dimension at most k (puncturing cannot increase dimension). For C' to have distinct codewords, we must have $|C'| \leq 2^{n-(d-1)}$. But $|C'| = |C| = 2^k$, so:

$$2^k \leq 2^{n-d+1} \implies k \leq n - d + 1.$$

□

Codes achieving equality, $d = n - k + 1$, are called *Maximum Distance Separable (MDS)* codes and are extremely rare and valuable. For example, Reed Solomon Codes (see next chapter) achieve it.

Remark 1.7.2. *Singleton Bound means one cannot simultaneously maximize information rate (k/n) and error correction capability (d).*

Theorem 1.7.3 (Hamming (Sphere-Packing) Bound). *If a code $C \subseteq \mathbb{F}_2^n$ corrects $t = \lfloor (d - 1)/2 \rfloor$ errors, then:*

$$\sum_{i=0}^t \binom{n}{i} \leq 2^{n-k}.$$

Proof. If C corrects t errors, the Hamming balls $B_t(c)$ of radius t around each codeword are pairwise disjoint (by the triangle inequality: if $d(x, c_1) \leq t$ and $d(x, c_2) \leq t$ with $c_1 \neq c_2$, then $d(c_1, c_2) \leq 2t < d$, contradiction). Each ball has size $\sum_{i=0}^t \binom{n}{i}$. Since $|C| = 2^k$ such balls fit inside $\mathbb{F}_2^n$:

$$2^k \cdot \sum_{i=0}^t \binom{n}{i} \leq 2^n.$$

□

Perfect codes achieve equality in this bound — they waste no space. They are extremely rare (Hamming codes, Golay codes, and trivial cases).

Upper bounds tell us the ceiling. But they don't guarantee anything lives up there. Enter Gilbert-Varshamov.

Theorem 1.7.4 (Gilbert-Varshamov Bound). *For any n and d , there exists a linear $[n, k, d]$ code with:*

$$k \geq n - \log_2 \sum_{i=0}^{d-2} \binom{n}{i}.$$

Proof. Construct a parity-check matrix H greedily. We need each nonzero vector in $\mathbb{F}_2^n$ to be expressible as a linear combination of at most $d - 1$ columns of H . Start with an empty H and add columns one by one. As long as the number of such linear combinations is less than 2^{n-k} (the number of possible syndromes), we can always find a column that hasn't appeared yet. The number of forbidden linear combinations is:

$$\sum_{i=0}^{d-2} \binom{n}{i}.$$

If $\sum_{i=0}^{d-2} \binom{n}{i} < 2^{n-k}$, we can always add a new column. This gives $k \geq n - \log_2 \sum_{i=0}^{d-2} \binom{n}{i}$. □

Remark 1.7.5. *This is a constructive existence result: it guarantees that for any choice of (n, d) , we can always build a code satisfying these parameters. Unlike the upper bounds, GV doesn't define an unattainable ideal — it shows what's always achievable.*

The open question

Upper bounds define what's impossible; GV shows what's always possible. The bounds are “relatively tight” — there's always a code somewhere in that feasible region, so we're not designing in a void.

But how tight are they really? GV tells us $k \geq \dots$, and Hamming tells us $k \leq \dots$. There is a gap between these two inequalities.

Can we construct codes that beat Gilbert-Varshamov while respecting the upper bounds? For specific families (like Hamming or BCH codes), yes — sometimes they outperform GV significantly. But in the general asymptotic regime, this gap remains largely open.

Closing this gap is one of the great problems in coding theory.

Chapter 2

Examples of Codes

2.1 Hamming Codes

2.1.1 The idea

We want a parity matrix H that tells us where an error has occurred. If e contains one non-zero bit at position i , $s = He^T$ will be the i -th column of H . In order to identify the column number, we construct H such that its i -th column is equal to i in binary. As H is an $m \times n$ matrix whose columns are all nonzero vectors, one can count at most $2^m - 1$ different column numbers. This is how we choose n : $n = 2^m - 1$, so that one column of H corresponds to one error position.

With this idea, a Hamming code is always $[2^m - 1, 2^m - m - 1]$. Finally, $d = 3$. If $w_H(c) = 1$ with 1 at position i , $Hc^T = \text{column } i$, which is nonzero. If $w_H(c) = 2$ with 1s at positions i and j , $Hc^T = \text{column } i \oplus \text{column } j$, which is nonzero since all columns are distinct. However, when $w_H(c) = 3$ with 1s at positions i, j, l such that

$$\text{col}_i \oplus \text{col}_j = \text{col}_l$$

all XOR to 0. This means that a codeword of weight 3 exists.

Then, it is $[2^m - 1, 2^m - m - 1, 3]$. We call them SEC (Single Error Correction).

Remark 2.1.1. *Hamming codes are perfect codes. They achieve the sphere-packing bound with equality.*

Proof. Consider a Hamming code $[2^m - 1, 2^m - m - 1, 3]$ with minimum distance $d = 3$, so it corrects $t = \lfloor (d - 1)/2 \rfloor = 1$ error. The sphere-packing bound states:

$$2^k \cdot \sum_{i=0}^t \binom{n}{i} \leq 2^n.$$

Substitute the Hamming parameters:

$$2^{2^m - m - 1} \cdot \sum_{i=0}^1 \binom{2^m - 1}{i} \leq 2^{2^m - 1}.$$

Evaluate the binomial sum:

$$\sum_{i=0}^1 \binom{2^m - 1}{i} = \binom{2^m - 1}{0} + \binom{2^m - 1}{1} = 1 + (2^m - 1) = 2^m.$$

Thus:

$$2^{2^m - m - 1} \cdot 2^m = 2^{2^m - m - 1 + m} = 2^{2^m - 1}.$$

The inequality becomes an equality, so Hamming codes are **perfect codes**: they achieve the maximum possible error-correction capability for their length and dimension. $\square$

2.1.2 Example: the $[7, 4]$ Hamming code

Here is an example with the Hamming code $[7, 4, 3]$:

$$H = \begin{pmatrix} 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ 0 & 1 & 1 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 & 1 & 1 & 1 \end{pmatrix}.$$

Look at H 's columns: $(001)_2 = 1, (010)_2 = 2, \dots, (111)_2 = 7$.

A generator matrix can be derived from H , and after Gaussian elimination, one systematic form is:

$$G = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 1 \end{pmatrix} = (I_4 \mid P),$$

where P is the 4×3 parity matrix. A message (m_1, m_2, m_3, m_4) is encoded as

$$c = (m_1, m_2, m_3, m_4, m_1 + m_3 + m_4, m_2 + m_3 + m_4, m_1 + m_2 + m_3 + m_4),$$

where the last three bits are the parity checks.

Let us give an example of encoding and decoding with Hamming $[7, 4]$ code:

- Codeword: $[1 \ 0 \ 1 \ 1 \ 0 \ 1 \ 0]$
- Received: $[1 \ 1 \ 1 \ 1 \ 0 \ 1 \ 1]$

There are two errors, at positions 2 and 7. The syndrome is $2 \oplus 7 = 5$. Therefore, the syndrome is $[1 \ 0 \ 1]$, which corresponds to position 5. The decoder assumes that there is one error at position 5 and flips that bit:

- Corrected word: $[1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1]$
- Decoded message: $[1 \ 1 \ 1 \ 1]$

However, this correction is wrong. The original errors at positions 2 and 7 are still present, and flipping position 5 introduces a third error. A standard Hamming $[7, 4]$ code can correct one error, but it cannot reliably detect two errors.

2.1.3 SECDED Codes

An easy way to prevent this problem when two errors occur is by adding a global parity bit: take $[2^m, 2^m - m]$ instead of $[2^m - 1, 2^m - m - 1]$ and write

$$H = \begin{pmatrix} 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 & 1 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 & 1 & 0 \end{pmatrix}$$

The matrix H has its first row full of 1s and columns permuted so G is in systematic form.

They are called SECDED (for Single Error Correction Double Error Detection). When $y = xG + e$ is received, the syndrome $s = Hy^T$ is then an m -dimensional vector and its first bit is the global parity check. Since the first row of H contains all 1s, this first bit is the XOR of all components in y . When y is a codeword c , we know that $Hc^T = 0$, meaning that there are no errors. This also implies that a valid codeword always has an even number of 1s. The global parity bit tells us exactly: the parity is OK if it equals 0, not OK otherwise. One error breaks the parity, while two errors restore it. These codes always assume that there are at most two errors, and cannot handle more.

Syndrome	Parity	Errors	Action
0	0	0	—
$\neq 0$	1	1	Corrects the one error
$\neq 0$	0	2	Detects 2 errors, no correction

Table 2.1: SECDED error correction behavior

Let's summarize:

See `src/hamming.py` for an implementation.

For instance, SECDED is pervasive in RAM and low-latency environments: it is computationally fast to decode and has negligible latency overhead, while remaining effective given that real-world memory errors rarely exceed two bits.

2.2 Reed-Solomon Codes

2.3 Low Density Parity Check (LDPC) Codes

Chapter 3

Error-Correcting Codes in Cryptography

3.1 Motivation

Modern Cryptography faces the incoming quantum computer that breaks asymmetric primitives like RSA, Diffie-Hellman, Elliptic Curves,... all used currently for key exchange and digital signatures. Error-Correcting Codes give a natural way to create some quantum-resistant key exchange cryptosystems.